


AI (LLM's) are everywhere!

- But how much Production Know-How do we have...?
- What are their TCO...?
- How to design a system given expected requirements?
 - Model type
 - Model size
 - # User requests
 - Response times
 - Etc

LLM Model Size Impact on Application Performance

Juan Jose Olivera Loyola



PPTM Final Project - Professor Jesus Jose Labarta
TMIRI-HPC

Goals

1. Understand intrinsic nature of hardware utilization of LLM applications.
 - a. Computation Structure, Compute vs Memory Boundness, Characteristic Values
2. Be able to derive characteristic numbers of LLM computation in llama.cpp
 - a. Formulas of token gen/sec, ratios, times, etc.
3. Have enough information to be able to compare other models or other applications/implementations

What is llama.cpp (the profiled application)?


Ollama: A LLM Containerization and Inference Platform

.... Docker for LLM's

llama.cpp: Inference engine behind Ollama

- Created in February 2023
- Spinoff of GGML a matrix multiplication library
- Creator of the common **GGUF model format**.
- Referenced even by Nvidia Technical Blog:

[Accelerating LLMs with llama.cpp on NVIDIA RTX Systems | NVIDIA Technical Blog](#)




LLM Models

Model Name (GGUFs)	# Layers	Quantization (avg. bits per weight)	(GB)
Llama 3B-Q4_K_M	28	4.8 bits	2GB
Qwen 7B-Q4_K_L	28	5.1 bits	5GB
Qwen 14B-Q4_K_L	48	5.1 bits	9.5GB
Qwen 32B-Q4_K_L	64	5.1 bits	20.4GB
Llama 70B-Q4_K_M	80	4.8 bits	42GB

Family of GPT (Generative Pre-Trained Transformer)
Models


Total Prompt Execution Time (1000 tokens)



AI and GPT's use heavily
Matrix Multiplications
 $\sim O(N^{2.37})$ to $O(N^3)$

...Seems to scale linearly
with 1 GPU

Project Tools & Methodology



Slurm Compilation Script

```
export SRUN_CPUS_PER_TASK=${SLURM_CPUS_PER_TASK}

module load cmake
module load nvidia-hpc-sdk/25.1
module load binutils

cd "$PWD/llama.cpp"
rm -rf build
cmake -DCMAKE_C_COMPILER=gcc -DLLAMA_AVX_VNNI=ON -DCMAKE_CXX_FLAGS="-march=native" -DCMAKE_CXX_COMPILER=g++ -B build
-DGGML_CUDA=ON -DCMAKE_CUDA_ARCHITECTURES="90;90;90;90"

cmake --build build --config Release -j 80
```

Slurm Profile Script

```
#!/bin/sh
#SBATCH --job-name=profile_llamacpp
#SBATCH --output=sjobs/profile_llamacpp_%j.out
--error=sjobs/profile_llamacpp_%j.error
#SBATCH --nodes=1
#SBATCH --ntasks-per-node=80
#SBATCH --constraint=perfparanoid
#SBATCH --qos=acc_debug
# Model loaded from llms/
modelFile="DeepSeek-R1-Distill-Llama-3B-Q4_K_M.gguf"                                # For skipping program setup
and avoid large trace useless traces
inputPrompt="Explain, in at least 800 words, the usage of HPC on addressing modern world problems and its impact on economy"
nTokens=1000
GPU
# CPU metrics/hw counters sampling frequency (Hz)
# GPU metrics/hw counters sampling frequency (Hz)
nsysFileName="run-$modelFile-$nTokens.nsys-rep"
baseFileName="run-$modelFile-$nTokens"
export SRUN_CPUS_PER_TASK=${SLURM_CPUS_PER_TASK}

module load cuda/12.6 nvidia-hpc-sdk/25.1

#SBATCH -D .
#SBATCH
#SBATCH --time=00:20:00
#SBATCH --ntasks=80
#SBATCH --cpus-per-task=1
#SBATCH --gres=gpu:4
#SBATCH -A nct_323
traceOffset=0
modelGpuLayers=10000
cpuSamplingFreq=231250
gpuSamplingFreq=10000
```


Slurm Profile Script

```
CUDA_VISIBLE_DEVICES=0 nsys profile \  
--output "experiments/$baseFileName" \  
--sample=process-tree --sampling-period=$cpuSamplingFreq \  
--trace=cuda,nvtx,osrt,cublas,cudnn --cuda-graph-trace=node \  
--gpu-metrics-devices=cuda-visible --gpu-metrics-frequency=$gpuSamplingFreq \  
--cpuctxsw=process-tree \  
--delay $traceOffset \  
./llama.cpp/build/bin/llama-cli -m "llms/$modelFile" -no-cnv \  
-n $nTokens -nlg $modelGpuLayers -p "$inputPrompt"  
  
cd experiments  
module load intel mkl python/3.12.1 sqlite3  
nsys2prv -t cuda_api_trace,nvtx_pushpop_trace,graph,gpu_metrics "$nsysFileName" "$baseFileName"  
tar -cvjf "$baseFileName.prv.tar.bz2" "$baseFileName.pcf" "$baseFileName.prv" "$baseFileName.row"  
rm "$baseFileName.pcf" "$baseFileName.prv" "$baseFileName.row" "$baseFileName.sqlite"
```

Project Tools & Methodology

1. Generate Traces	2. Analysis of 32B Model (Identifying Structure)	3. Analysis of rest of Models (Focus of Analysis)	4. Characterizing Performance (Perf Model & KPIs)
--------------------	---	--	--

Application Work excludes Initialization



Initialization/Model Loading

~ 35% total time for 1 1000 token prompt Llama-3B

Calls: CUDA profiling init, fopen, mmap, ioctl, cudaMalloc/cpy.

Real Work: Token Prediction


Calls: Kernels, CUDA Sync,
memCpy, etc.

Identifying Structure

Setup Work
(Repetitive) Body Work
Termination Work

Model: **Qwen-32B-Q4_K_L**
-n Tokens = 1000
-p “Explain the usage of HPC on addressing modern world problems and its impact on economy” (20 - 25 tokens)

Llama Inference Process




How? By finding Token Generation Structure

How? By finding Token Generation Structure

cuGraphLaunch Calls = 999 + 1 Graph Instantiation = 1000 Tokens

How? By finding Token Generation Structure


cuGraphLaunch Calls = 999 + 1 Graph Instantiation = 1000 Tokens




Zoom into 1 Iteration

63 normal sub iterations + 1 PostWork Subiteration

x64 soft_max_f32 = # layers in Qwen2.5 32B




Zoom into 1 Sub Iteration




- End
- cpy_f32_f16
- k_bin_bcast
- mul_mat_vec
- mul_mat_vec_q
- quantize_q8_1
- rms_norm_f32
- soft_max_f32
- unary_op_kernel
- rope_neox


Zoom into 1 Sub Iteration




Zoom into 1 Sub Iteration



Zoom into 1 Sub Iteration



Kernels to LLM Architecture



2x rms_norm_f32

2x rope_neox

1x mul_mat_vec_q?

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$


1x soft_max_f32 2x (mul_mat_vec,
 mul_mat_vec_q)

LLaMA

$$\text{FFN}_{\text{SwiGLU}}(x, W, V, W_2) = (\text{Swish}_1(xW) \otimes xV)W_2$$

We use the swish function with $\beta = 1$. In this case it's called the **Sigmoid Linear Unit (SiLU)** function.


$$\text{swish}(x) = x \text{ sigmoid}(\beta x) = \frac{x}{1 + e^{-\beta x}}$$




1x unary_op_kernel

3x mat_mul_vec_q

Zoom into 1 Sub Iteration




Start and Ending of Iteration




Start: cudaMemcpyHostToDevice

End: Bigger mul_mat_vec_q, cudaMemcpyDeviceToHost

Final Subiteration is Last FF




Iteration Times ~ Time Between cudaGraphLaunch()




Time Between cudaGraphLaunch() in all trace

Time Between cudaGraphLaunch @ ru1-DeepSeek-R1-Distill-Qwen-32B-Q4_K_L.gguf-1000.juntar.prv




How many big times between cudaGraphLaunch?




34 long calls / 1000 tokens = 3.4% or 1 every 30 tokens

Time Between cudaGraphLaunch @ ru1-DeepSeek-R1-Distill-Qwen-32B-Q4_K_L.gguf-1000.juntar.prv




Regular CPU - GPU Synchronization



Long CPU - GPU Synchronization

	THREAD 1.1.1
cuKernelGetFunction	0.07 %
cuLaunchKernel	0.49 %
cudaGraphDestroy_v10000	2.22 %
cudaGraphExecUpdate_v10020	2.21 %
cudaGraphGetNodes_v10000	0.11 %
cudaGraphKernelNodeGetParams_v10000	2.89 %
cudaGraphLaunch_v10000	10.56 %
cudaGraphNodeGetType_v10000	1.89 %
cudaLaunchKernel	10.94 %
cudaMemcpyAsync	0.46 %
cudaStreamBeginCapture_v10000	0.02 %
cudaStreamEndCapture_v10000	1.74 %
cudaStreamSynchronize	66.40 %

CUDA call @ run-DeepSeek-R1-Distill-Qwen-32B-Q4_K_L.gguf-1000.juntar.chop1.prv



Identifying Structure Recap


- Token generation = 1 iteration identified by cudaGraphLaunch()
- # iterations = **# tokens**
- Each iteration has **#layers** subiterations
- 30% of subiteration time is Multi-Head Attention
- 70% of subiteration time is FeedForward NN
- All iterations begin and end with hostToDeviceCpy and deviceToHostCpy
- The final subiteration has a larger matrix multiplication
- Some iterations (3.4%) require bigger setup time for updating cuda graph parameters

Focus of Analysis (FoA)


What impacts the most?

What changes most when varying model size?

Kernel Percentage



Kernel Relative Time Increase




MHA vs FF Work Proportion


Zoom into 1 Sub Iteration



Slide 24 Identifying Structure for Qwen 32B-Q4_K_M




MHA vs FF Work Proportion



mul_mat_vec_q
dominates the time

Subiteration Kernels Proportion

Llama-3B



Qwen-7B



Qwen-14B



Qwen-32B




Llama-70B



% Stream Multiprocessors (SM) Active


H100 SMs: 114



Render Config	
Shading Units:	14592
TMUs:	456
ROPs:	24
SM Count:	114
Tensor Cores:	456
L1 Cache:	256 KB (per SM)
L2 Cache:	50 MB

Bandwidth Throughput (%)

H100 BW: ???



Common Opinion that Memory is Bottleneck

LLaMA explained: KV-Cache, Rotary Positional Embedding, RMS Norm, Grouped Query Att... ⏱ ➡ ⓘ

GPUs have a “problem”: they’re too fast.


- In recent years, GPUs have become very fast at performing calculations, insomuch that the speed of computation (FLOPs) is much higher than the memory bandwidth (GB/s) or speed of data transfer between memory areas. For example, an NVIDIA A100 can perform 19.5 TFLOPs while having a memory bandwidth of 2TB/s.
- This means that sometimes the bottleneck is not how many operations we perform, but how much data transfer our operations need, and that depends on the size and the quantity of the tensors involved in our calculations.
- For example, computing the same operation on the same tensor N times may be faster than computing the same operation on N different tensors, even if they have the same size, this is because the GPU may need to move the tensors around.
- This means that our goal should not only be to optimize the number of operations we do, but also minimize the memory access/transfers that we perform.**

NVIDIA A100 TENSOR CORE GPU SPECIFICATIONS (SXM4 AND PCIe FORM FACTORS)			
A100	A100	A100	A100
40GB PCIe	80GB PCIe	40GB SXM	80GB SXM
FP64	9.7 TFLOPS		
FP64 Tensor Core	19.5 TFLOPS		
FP32	19.5 TFLOPS*		
Tensor Float 32 (TF32)	154 TFLOPS 1,312 TFLOPS*		
BF16 QAT16 Tensor Core	312 TFLOPS 424 TFLOPS*		
FP16 Tensor Core	312 TFLOPS 424 TFLOPS*		
INT8 Tensor Core	624 TOPS 1,248 TOPS*		
GPU Memory	4GB HBM2	8GB HBM2	4GB HBM2
GPU Memory Bandwidth	1,555GB/s	1,935GB/s	1,555GB/s
Max Thermal Design Power (TDP)	250W	300W	400W
Multi-Instance GPU	Up to 7 MiGs @	Up to 7 MiGs @	Up to 7 MiGs @

*Actually 40 times slower, because each operation is 32 bit.


* 1555 GB/s via HBM A100 server boards; PCIe 64GB/s via NVLink Bridge for up to two GPUs

NO Tensor Cores Active per Iteration



0% Tensor Cores Utilization

NO Tensor Cores Active per Iteration



nsys-ui also shows no Tensor Activity

0% Tensor Cores Utilization

=> Everything is in Traditional CUDA Threads

Focus of Analysis Recap

- **mul_mat_vec_q \propto Model Size.** FF mostly has mul_mat_vec_q
- **mul_mat_vec, k_bin_bcast \propto sqrt(Model Size).**
- FF grows bigger than MHA
- The Memory Bandwidth is higher for MHA portion and Lower in End of FF
- No Tensor Cores Utilization so all computations are in regular SM Warps/Threads
- **FF is the most Compute Heavy, so mul_mat_vec_q is more impacted by compute bottlenecks than memory in llama.cpp with H100**


Thank You

Appendix

References


- bartowski. (2025) Models GGUFs [bartowski/DeepSeek-R1-Distill-Llama-70B-GGUF · Hugging Face](#), [bartowski/DeepSeek-R1-Distill-Qwen-32B-GGUF · Hugging Face](#), [bartowski/DeepSeek-R1-Distill-Qwen-14B-GGUF · Hugging Face](#), [bartowski/DeepSeek-R1-Distill-Qwen-7B-GGUF · Hugging Face](#), [hassenhamdi/DeepSeek-R1-Distill-Llama-3B-GGUF · Hugging Face](#)
- Chockalingam, A (2024). [Accelerating LLMs with llama.cpp on NVIDIA RTX Systems | NVIDIA Technical Blog](#)
- QwenTeam, Alibaba Group. QWEN2 TECHNICAL REPORT. [2407.10671](#) Pag 5
- Meta. The Llama3 Herd of Models [2407.21783](#) Pag 7.
- Umar Jamil (2024) LLaMA explained: KV-Cache, Rotary Positional Embedding, RMS Norm, Grouped Query Attention, SwiGLU. [\(132\) LLaMA explained: KV-Cache, Rotary Positional Embedding, RMS Norm, Grouped Query Attention, SwiGLU - YouTube](#)
- Kansal, V. (2023) Understanding the GGUF Format: A Comprehensive Guide. [Understanding the GGUF Format: A Comprehensive Guide | by Vimal Kansal | Medium](#)
- Quantization Bits Table. [GGUF quantizations overview](#)

Setup Execution + First Token Iteration



2 Irregular Execution iterations happen before GraphLaunched (blue calls)
iterations. Irregular execution uses more kernel types

Something odd... There are 2 Graph Instantiations...



Kernels by Iteration Type


Profile 1 (Early) Iteration

cpy_f32_f16	128
k_bin_bcast	513
k_get_rows_float	2
mul_mat_vec	128
mul_mat_vec_q	449
quantize_q8_1	449
rms_norm_f32	129
soft_max_f32	64
unary_op_kernel	64
rope_neox	128


Profile 2 (Late) Iteration

Kernel2		
convert_unary	64	← New
cpy_f32_f16	128	← New
k_bin_bcast	513	
k_compute_batched_ptrs	64	← New
k_get_rows_float	2	
mul_mat_vec	64	
mul_mat_vec_q	449	← Reduced
quantize_q8_1	449	
rms_norm_f32	129	
soft_max_f32	64	
unary_op_kernel	64	
rope_neox	128	

SubIteration Profile 2



SubIteration Profile 2



Marenostrum 5 profile llama.cpp model

```
CUDA_VISIBLE_DEVICES=0,1,2,3 nsys profile --output  
experiments/longrun_job_12_delay_all_duration_graph_mode.nsys-rep  
--sample=process-tree --sampling-period=1850000  
--trace=cuda,nvtx,osrt,cublas,cudnn --cuda-graph-trace=graph  
--gpu-metrics-devices=cuda-visible --gpu-metrics-frequency=10000  
--cpuctxsw=process-tree --delay 12 ./llama.cpp/build/bin/llama-cli -m  
LLM_Models/DeepSeek-R1-Distill-Qwen-32B-Q4_K_L.gguf -no-cnv -n 1000 -ngl  
1000 -p "Explain the usage of HPC on addressing modern world problems and its  
impact on economy"
```

Marenostrum 5 convert nsys to paraver

```
module load intel mkl python/3.12.1 sqlite3  
  
nsys2prv -t cuda_api_trace,nvtx_pushpop_trace,graph,gpu_metrics  
longrun_job_12_delay_all_duration_graph_mode.nsys-rep  
longrun_job_12_delay_all_duration_graph_mode  
  
tar -cvjf longrun_job_12_delay_all_duration_graph_mode.prv.tar.bz2  
longrun_job_12_delay_all_duration_graph_mode.pcf  
longrun_job_12_delay_all_duration_graph_mode.prv  
longrun_job_12_delay_all_duration_graph_mode.row
```